

The Tabular Data Model and SQL



Outline

- What Grammar of Data do tables follow? How is this grammar implemented in Pandas and SQL?
- Creation and alteration in the grammar of data
- The grammar as related to combining tables

A simple but powerful model of a table

- A collection of tables related to each other through common data values.
- Rows represent attributes of something
- Everything in a column is values of *one* attributes
- A cell is expected to be atomic
- Tables are related to each other if they have columns called keys which represent the same values

Contributors

Table:  

	id	last_name	first_name	middle_name	street_1	street_2	city	state	zip	amount	date	candidate_id
	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter
1	1	Agee	Steven	NULL	549 Laurel ...	NULL	Floyd	VA	24091	500	2007-06-30	16
2	5	Akin	Charles	NULL	10187 Suga...	NULL	Bentonville	AR	72712	100	2007-06-16	16
3	6	Akin	Mike	NULL	181 Baywo...	NULL	Monticello	AR	71655	1500	2007-05-18	16
4	7	Akin	Rebecca	NULL	181 Baywo...	NULL	Monticello	AR	71655	500	2007-05-18	16
5	8	Aldridge	Brittni	NULL	808 Capitol...	NULL	Washington	DC	20024	250	2007-06-06	16
6	9	Allen	John D.	NULL	1052 Cann...	NULL	North Augu...	SC	29860	1000	2007-06-11	16
7	10	Allen	John D.	NULL	1052 Cann...	NULL	North Augu...	SC	29860	1300	2007-06-29	16
8	11	Allison	John W.	NULL	P.O. Box 10...	NULL	Conway	AR	72033	1000	2007-05-18	16
9	12	Allison	Rebecca	NULL	3206 Sum...	NULL	Little Rock	AR	72227	1000	2007-04-25	16

Typing our tables

```
DROP TABLE IF EXISTS "candidates";  
DROP TABLE IF EXISTS "contributors";  
  
CREATE TABLE "candidates" (  
    "id" INTEGER PRIMARY KEY NOT NULL ,  
    "first_name" VARCHAR,  
    "last_name" VARCHAR,  
    "middle_name" VARCHAR,  
    "party" VARCHAR NOT NULL  
);
```

```
CREATE TABLE "contributors" (  
    "id" INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,  
    "last_name" VARCHAR,  
    "first_name" VARCHAR,  
    "middle_name" VARCHAR,  
    "street_1" VARCHAR,  
    "street_2" VARCHAR,  
    "city" VARCHAR,  
    "state" VARCHAR,  
    "zip" VARCHAR,  
    "amount" INTEGER,  
    "date" DATETIME,  
    "candidate_id" INTEGER NOT NULL,  
    FOREIGN KEY(candidate_id) REFERENCES candidates(id)  
);
```

Grammar of Data

Formalized by Hadley Wickham in `dplyr`¹.

1. provide simple verbs for simple things. These are functions corresponding to common data manipulation tasks
2. the backend does not matter. Here we constrain ourselves to Pandas and SQL in `sqlite`
3. multiple backends implemented in Pandas, Spark, Impala, Pig, `dplyr`, `ibis`, `blaze`

¹ Hadley Wickham: <https://cran.rstudio.com/web/packages/dplyr/vignettes/introduction.html>

Why bother with SQL and the grammar

- learn how to do core data manipulations, no matter what the system
- relational databases critical for tables that don't fit in memory.
- SQL is declarative, the query optimizer decides how to make the query.
- dbs like Cloudera, Azure Data Lake implement the SQL lingua-franca or similar over clusters, multiple databases

You can always find out what these core verbs are in any system!

VERB	dplyr	pandas	SQL
QUERY/SELECTION	filter() (and slice())	query() (and loc[], iloc[])	SELECT WHERE
SORT	arrange()	sort_values()	ORDER BY
SELECT-COLUMNS/ PROJECTION	select() (and rename())	[](__getitem__) (and rename())	SELECT COLUMN
SELECT-DISTINCT	distinct()	unique(),drop_duplicates()	SELECT DISTINCT COLUMN
ASSIGN	mutate() (and transmute())	assign	ALTER/UPDATE
AGGREGATE	summarise()	describe(),mean(),max()	None, AVG(),MAX()
SAMPLE	sample_n() and sample_frac()	sample()	implementation dep, use RAND()
GROUP-AGG	group_by/summarize	groupby/agg, count, mean	GROUP BY
DELETE	?	drop/masking	DELETE/WHERE

SQLITE

Sqlite is a on-file database, as opposed to other common databases such as Oracle and Postgres, which run as different processes on your system. Sqlite is great for on-disk large databases which wont fit into memory.

Its also built into Python, but to use the command line tool, I recommend you install it: <https://www.sqlite.org/download.html>. I also recommend you download and install the sqlite browser: <http://sqlitebrowser.org> .

Python implements a standard database API over all databases. Its called DBAPI2. It works across many SQL databases, including Sqlite. There is an even higher level API available, called [SQLAlchemy](#).

Connect to and populate the database

```
def get_db(dbfile):
    # get a connection
    sqlite_db = sq3.connect(dbfile)
    return sqlite_db

def init_db(dbfile, schema):
    """Creates the database tables."""
    db = get_db(dbfile)
    # execute some SQL code
    c = db.cursor()
    c.executescript(schema)
    # make a commit
    db.commit()
    return db
```

```
db=init_db("/tmp/cancont.db", schema)
# populate the database
dfcand.to_sql("candidates", db,
             if_exists="append", index=False)
dfcwc.to_sql("contributors", db,
            if_exists="append", index=False)
```

Using the [sqlite dbapi2](#):

```
# get a pointer to rows of results
c=db.cursor()
# you can also use db directly
c.execute("SELECT * FROM candidates;")
# you can fetch all, or some at a time
# here we fetch all
c.fetchall()
```

```
[(16, 'Mike', 'Huckabee', None, 'R'),
 (20, 'Barack', 'Obama', None, 'D'),
 (22, 'Rudolph', 'Giuliani', None, 'R'),
 (24, 'Mike', 'Gravel', None, 'D'),
 (26, 'John', 'Edwards', None, 'D'),
 (29, 'Bill', 'Richardson', None, 'D'),
 ....
 (41, 'Fred', 'Thompson', 'D.', 'R')]
```

Operation	Pandas	SQL
QUERY	<code>dfcwci.query("state=='VA' & amount < 400")</code>	<code>SELECT * FROM contributors WHERE state='VA' AND amount < 400;</code>
SORT	<code>dfcwci.sort_values("amount")</code>	<code>SELECT * FROM contributors ORDER BY amount;</code>
SELECT-COLUMNS	<code>dfcwci[['first_name', 'amount']]</code>	<code>SELECT first_name, amount FROM contributors;</code>
SELECT-DISTINCT	<code>dfcwci[['last_name', 'first_name']].drop_duplicates()</code>	<code>SELECT DISTINCT last_name, first_name FROM contributors;</code>
ASSIGN	<code>dfcwci['name']=dfcwci['last_name']+", "+dfcwci['first_name']</code>	<code>ALTER TABLE contributors ADD COLUMN name; UPDATE contributors SET name = ? WHERE id = ?;</code>
AGGREGATE	<code>dfcwci.amount.max()</code>	<code>SELECT MAX(amount) FROM contributors;</code>
GROUP-AGG	<code>dfcwci.groupby("state").sum()</code>	<code>SELECT state,SUM(amount) FROM contributors GROUP BY state;</code>
DELETE	<code>dfcwci=dfcwci[dfcwci.last_name!='Ahrens']</code>	<code>DELETE FROM contributors WHERE last_name="Ahrens";</code>
DELETE TABLE	<code>del dfcwci</code>	<code>DELETE FROM contributors;</code>
CREATE TABLE	<code>dfcwci=pd.read_csv("data/contributors_with_candidate_id.txt")</code>	<code>INSERT INTO candidates (id, first_name, last_name, middle_name, party) VALUES (?, ?, ?, ?, ?);</code>

Some examples of the verbs

combining multiple selects

```
dfcwci[(dfcwci.state=='VA') & (dfcwci.amount < 400)]
```

```
SELECT * FROM contributors WHERE state='VA' AND amount < 400;
```

array membership

```
dfcwci[dfcwci.state.isin(['VA', 'WA'])]
```

```
SELECT * FROM contributors WHERE state IN ('VA', 'WA');
```

project onto specific columns

```
dfcwci[['first_name', 'amount']]
```

```
SELECT first_name, amount FROM contributors;
```

aggregation

```
dfcwci.amount.max()
```

```
SELECT MAX(amount) FROM contributors;
```

The Split-Applv-Combine (GROUP-AGG) pattern

1. split the data into groups
2. based on some criteria apply a function to each group independently
3. combine the results into a data structure

```
SELECT state, MEAN(amount)
FROM contributors GROUP BY state;
```

```
dfcwci.groupby("state")["amount"].mean()
```

```
state
AK      403.333333
AR     1183.333333
AZ      120.000000
CA     -217.988261
CO    -1455.750000
CT     2300.000000
DC     -309.982000
...
```

Creation and Alteration

Inserting data into SQLITE

```
sql_template = """
INSERT INTO candidates (id, first_name, last_name, middle_name, party) \
    VALUES (?, ?, ?, ?, ?);
"""

with open("data/candidates.txt") as fd:
    slines = [l.strip().split('|') for l in fd.readlines()]
    for line in slines[1:]:
        theid, first_name, last_name, middle_name, party = line
        valstoinsert = (int(theid), first_name, last_name, middle_name, party)
        db.cursor().execute(sql_template, valstoinsert)

db.commit()
```

What is this `commit()`?

Things can go wrong. Software fails. Hardware fails. Networks fail. Multiple clients may update tables.

You don't want a bank to only do a withdrawal but not a deposit on a money transfer! You want a **transaction**.

A transaction groups reads and writes together into ONE operation. Either all happen and succeed, or none do. The database will **rollback** any partial changes. `db.commit()` signals a transaction.

The acronym ACID was coined in 1983 by Theo Härder and Andreas Reuter to describe transactions:

- **A** : Atomic, the above all or none guarantee
- **C** : Consistency means “good state”. In bank transfer, a half done transfer is not consistent. This definition is application specific.
- **I** : Isolation means that concurrently executing transactions don't interfere with each other: one happens first, and the other only sees things after the first transaction executed
- **D** : Durability : after commit, data won't be lost.

Altering tables by adding columns

```
dfcwci['name']=dfcwci['last_name']+", "+dfcwci['first_name'] # pandas
```

```
alt="ALTER TABLE contributors ADD COLUMN name;" # SQL
```

```
db.cursor().execute(alt)
```

```
db.commit()
```

```
alt2="UPDATE contributors SET name = ? WHERE id = ?;" # SQL
```

```
for ele in out2:
```

```
    db.cursor().execute(alt2, ele)
```

```
db.commit()
```

Updating existing columns

```
dfcwci.loc[dfcwci.state=='VA', 'name']="junk" # pandas
```

```
upd="UPDATE contributors SET name = 'junk' WHERE state = 'VA';"  
db.cursor().execute(upd) # SQL  
db.commit()
```

Combining Tables

RELATIONSHIPS

- we usually need to combine data from multiple sources
- different systems have different ways, but most copy SQL
- Example: sub-select:


```
obamaid=dfcand.query("last_name=='Obama'")['id'].values[0]  
dfcwci.query("candidate_id==%i" % obamaid)
```

```
SELECT * FROM contributors WHERE  
    candidate_id = (SELECT id from candidates WHERE last_name = 'Obama');
```

Two table grammar: JOINS

MUTATING JOINS: add new variables to one table from matching rows in another. Types of join are:

1. inner join,
2. left(outer) join
3. right(outer) join, and
4. full(outer) join



INNER JOINS

Combine 2 tables on a common key value. 90% of the time this is an explicit inner join

left					right					Result						
	key1	key2	A	B		key1	key2	C	D		key1	key2	A	B	C	D
0	K0	K0	A0	B0	0	K0	K0	C0	D0	0	K0	K0	A0	B0	C0	D0
1	K0	K1	A1	B1	1	K1	K0	C1	D1	1	K1	K0	A2	B2	C1	D1
2	K1	K0	A2	B2	2	K1	K0	C2	D2	2	K1	K0	A2	B2	C2	D2
3	K2	K1	A3	B3	3	K2	K0	C3	D3							

PANDAS:

```
dfcwc.merge(dfccand, left_on="candidate_id", right_on="id")
```

SQL:

```
SELECT * FROM
    contributors JOIN candidates
ON contributors.candidate_id = candidates.id;
```

Why are these joins important?

Denormalized data is very important for data analysis. Why?

Foreign keys carry no semantic meaning and models need to be fed data rather than pointers.

If we have many tables with information about items, we want to join them rather than keep the pointers.

This wide, denormalized form can be fed to models.

Close the database connection

```
db.close()
```